

1 **Part 4: Understanding Data Sets**

2 **Summary**

3 This Part presents an example of reading in a data set, and ensuring the
4 data is fully understood and properly represented prior to any further
5 analysis.

6 The ability to read in, clean, and understand data sets is a crucial skill.
7 Data do not typically come in a convenient form. Careful consideration
8 should be given to exploring the data to ensure full understanding of the
9 nature of the data set.

1 **Example Data Set**

2 First, we will read in a data set that we will use in our following examples.

3 Visit the website

4 https://www.sec.gov/opa/data/market-structure/marketstructuredownloadshhtml-by_security.html

5 and download the data from the first quarter of 2017. Create a directory
6 for working on these examples, and place this file in that directory.

7 These data come from the U.S. Securities and Exchange Commission re-
8 spository on market structure. This data set contains daily characteristics
9 of a range of market variables over 5000 equities and ETFs, accumulated
10 over several exchanges.

- 1 Be sure that the working directory is correctly set in R.
- 2 Now use the `read.table()` function to read in the file:

```
> fulldata = read.table("q1_2017_all.csv", sep=",",  
+                        header=T, quote="")
```

- 3 This is a “csv” file, i.e., there are “comma-separated-values”. This leads to
- 4 the use of `sep=", "`. The `header=T` argument specifies that the first line
- 5 of the file gives the variable names.

1 **Exercise:** What happens if you use `read.table()` in the example above
2 without `quote=" "`?

3 _____

4 _____

5 _____

6 _____

7 _____

8 _____

9 _____

10 _____

11 _____

12

1 Understanding the Variables

2 Whenever a new data set is encountered, one should fully understand its
3 variables. Specific questions to address for each variable:

4 1. Describe the information that is stored in this column.

5 2. What type of variable is it? (Timestamp, ordinal factor, categorical fac-
6 tor, ratio scale, other?) Be sure that R represents the variable correctly.

7 3. Are there any missing values? What is the source of the missingness?

8 4. Are there any extreme/inappropriate values? What is the source of
9 the extreme value?

1 Look at the variables:

```
> names(fulldata)
[1] "Date"                "Security"
[3] "Ticker"              "McapRank"
[5] "TurnRank"            "VolatilityRank"
[7] "PriceRank"           "LitVol..000."
[9] "OrderVol..000."      "Hidden"
[11] "TradesForHidden"     "HiddenVol..000."
[13] "TradeVolForHidden..000." "Cancels"
[15] "LitTrades"           "OddLots"
[17] "TradesForOddLots"    "OddLotVol..000."
[19] "TradeVolForOddLots..000."
```

- 2 Background on the data can be found in the associated README file, avail-
3 able on Canvas in the “Data Sets” folder in the “Files” page. We will refer-
4 ence portions of it below.

1 Column 1: Date

2 The first column in `fulldata` is a date stamp for the observation. R has
3 special functions for handling dates which will prove useful later when
4 working with time series.

5 The `as.Date()` function transforms into the date format:

```
> print(fulldata$Date[1])  # This is Jan. 3, 2017  
[1] 20170103
```

```
> fulldata$Date = as.Date(as.character(fulldata$Date),  
+                          format="%Y%m%d")
```

6 Note that it is necessary to cast the dates as strings, and then specify the
7 format of those strings.

1 **Exercise:** Explain the use of the `format` argument to `as.Date()`.

2

3

4

5

6

7

8

9

10

11

12

1 Columns 2 and 3: Security and Ticker

- 2 The variables `Security` and `Ticker` are already appropriately treated as
3 **factors** by R:

```
> is.factor(fulldata$Sec) & is.factor(fulldata$Tick)
[1] TRUE
```

- 4 We also note that `Security` is either `Stock` or `ETF` (Exchange Traded
5 Fund), and that there are over 5,000 equities under consideration:

```
> levels(fulldata$Security)
[1] "ETF" "Stock"
```

```
> nlevels(fulldata$Ticker)
[1] 5338
```

1 **Exercise:** How many of the different ticker symbols are ETFs?

2

3

4

5

6

7

8

9

- 1 **Exercise:** Execute the commands below, and discuss what you find:

```
> length(unique(fulldata$Date))  
> table(table(fulldata$Ticker))  
> which.max(table(fulldata$Ticker))  
> fulldata[duplicated(fulldata[,c(1,3)]),]
```

2

3

4

5

6

7

8

- 1 We will remove the duplicated lines using the following command:

```
> fulldata = fulldata[!duplicated(fulldata[,c(1,3)],  
+                               fromLast=TRUE), ]
```

- 2 Note that by using `fromLast=TRUE`, we remove the **first** of each pair of
- 3 duplicates.

1 **Columns 4 through 7: Rank Variables**

2 The next four columns provide ranks of the stocks/ETFs with respect to
3 key attributes of Market Capitalization, Turnover, Volatility, and Price.

4 In general, we would prefer to start with data in its most “raw” format,
5 and one could construct variables such as these from that data. Alas, we
6 are often left to work with what is available.

1 **Exercise:** Inspect the output of

```
> table(fulldata$VolatilityRank, fulldata$Security,  
+       fulldata$Date)
```

2 What does this tell us about the structure of these variables? Why should
3 we be very careful with these variables?

4 _____

5 _____

6 _____

7 _____

8 _____

9 _____

10

- 1 We will change these into **ordered factors** to reflect their ordinal nature:

```
> fulldata$McapRank =  
+   factor(fulldata$McapRank, ordered=TRUE)  
> fulldata$TurnRank =  
+   factor(fulldata$TurnRank, ordered=TRUE)  
> fulldata$VolatilityRank =  
+   factor(fulldata$VolatilityRank, ordered=TRUE)  
> fulldata$PriceRank =  
+   factor(fulldata$PriceRank, ordered=TRUE)
```

1 **Columns 8 thru 19: The Count Variables**

2 The remaining column in the data set consist of trade and share counts of
3 different types.

4 Note that some of the volume counts are reported in 1000's of shares, hence
5 there are non-integer values among the counts.

- 1 The `summary()` function is a useful first step in understanding the basic
- 2 properties of the variables. A primary concern is the presence of extreme
- 3 or impossible values.

```
> summary(fulldata[,8:19])
```

LitVol..000.	OrderVol..000.	Hidden	TradesForHidden
Min. : 0.00	Min. : 0	Min. : -143.0	Min. : 0
1st Qu.: 3.46	1st Qu.: 885	1st Qu.: 10.0	1st Qu.: 41
Median : 44.57	Median : 4003	Median : 75.0	Median : 501
Mean : 431.76	Mean : 35747	Mean : 454.4	Mean : 3476
3rd Qu.: 254.41	3rd Qu.: 15581	3rd Qu.: 370.0	3rd Qu.: 2846
Max. : 119014.06	Max. : 7293901	Max. : 89335.0	Max. : 533874

HiddenVol..000.	TradeVolForHidden..000.	Cancels
Min. : -7903.64	Min. : 0.00	Min. : 0
1st Qu.: 1.15	1st Qu.: 5.64	1st Qu.: 3126
Median : 8.74	Median : 55.31	Median : 15583
Mean : 66.91	Mean : 498.04	Mean : 73579
3rd Qu.: 44.02	3rd Qu.: 301.58	3rd Qu.: 60112
Max. : 14821.94	Max. : 133714.59	Max. : 7607714

LitTrades		OddLots		TradesForOddLots		OddLotVol..000.	
Min.	: 0	Min.	: 0	Min.	: 0	Min.	: 0.000
1st Qu.:	23	1st Qu.:	8	1st Qu.:	39	1st Qu.:	0.280
Median :	377	Median :	154	Median :	464	Median :	5.254
Mean :	2708	Mean :	916	Mean :	3115	Mean :	32.492
3rd Qu.:	2186	3rd Qu.:	887	3rd Qu.:	2558	3rd Qu.:	31.083
Max.	:383222	Max.	:130106	Max.	:462414	Max.	:3728.885
TradeVolForOddLots..000.							
Min.	: 0.00						
1st Qu.:	5.31						
Median :	49.94						
Mean :	422.78						
3rd Qu.:	260.18						
Max.	:108034.79						

- 1 **Exercise:** Do any of these variables take values that would seem to be
2 impossible? Can you speculate as to the source of these values?

3 _____

4 _____

5 _____

6 _____

7 _____

8 _____

9 _____


10 _____

11

- 1 Of course, graphical tools are useful for understanding the properties of
2 variables. This topic will be covered more fully in the next Part, but here
3 we consider a classic approach: the **boxplot**:

```
> boxplot(log(fulldata$Hidden),  
+         ylab="Log Count of Hidden Trades")
```

- 4 The “box” of the plot extends from the 25th to the 75th percentile of the
5 data, while the solid line in the middle of the box is at the median. The
6 two “arms” extend to the minimum and maximum value, **except** that any
7 value labelled an **outlier** is plotted separately.



- 1 **Exercise:** Why was the logarithm of the variable utilized in this case? What
2 issues did that create?

3 _____

4 _____

5 _____

6 _____

7 _____

8 _____

9 _____

10 _____

11 _____

12 _____

- 1 **Exercise:** Consider the outlier in the previous plot. Is this observation also
2 extreme in the other variables? Can you find the source of this behavior?

3 _____

4 _____

5 _____

6 _____

7 _____

8 _____

9 _____

10 _____

1 **Exercise:** Comment on the output of

```
> apply(fulldata[, 8:19], 2, which.max)
```

2 _____

3 _____

4 _____

5 _____

6 _____

7 _____

8 _____

9 _____

10 _____

11 _____

1 Missing Data

2 Most data sets contain some amount of missing data, for many reasons.

3 R uses the special symbol `NA` to represent a missing value.

4 As a first step, it is important to recognize how your data set repre-
5 sents missing values, so that R handles them properly. The function
6 `read.table()` has an argument `na.strings` to allow one to specify
7 missing values. **Be careful:** It is not unusual for data sets to use numbers
8 such as `-999` to indicate a missing value.

9 In a `.csv` file such as that we are using here, it is common to simply have
10 blank values to indicate missingness, i.e., there are consecutive commas
11 without a value between.

- 1 **Exercise:** Does R appropriately handle cases where a missing value is in-
- 2 dicated using blank values in a `.csv` file?

3 _____

4 _____

5 _____

6 _____

7 _____

8 _____

9 _____

10 _____

11 _____

12 _____

- 1 A useful tool is the function `md.pattern()` that is part of package `mice`.
- 2 This function shows the different “patterns” of missingness in the data set,
- 3 and how often they occur.

- 4 Consider the example on the following page, the result of calling

```
> md.pattern(data.frame(fulldata), plot=FALSE)
```

- 5 This shows that there are 325,082 complete cases in the data set. There is
- 6 one case where `VolatilityRank` is missing, and there are 76 cases where
- 7 `McapRank`, `TurnRank` and `PriceRank` are missing.

```

1      Date Security Ticker LitVol..000. OrderVol..000. Hidden
2 325082      1      1      1      1      1      1
3 76      1      1      1      1      1      1
4 1      1      1      1      1      1      1
5      0      0      0      0      0      0
6      TradesForHidden HiddenVol..000. TradeVolForHidden..000. Cancels
7 325082      1      1      1      1
8 76      1      1      1      1
9 1      1      1      1      1
10      0      0      0      0
11      LitTrades OddLots TradesForOddLots OddLotVol..000.
12 325082      1      1      1      1
13 76      1      1      1      1
14 1      1      1      1      1
15      0      0      0      0
16      TradeVolForOddLots..000. VolatilityRank McapRank TurnRank PriceRank
17 325082      1      1      1      1      1
18 76      1      1      0      0      0
19 1      1      0      1      1      1
20      0      1      76      76      76
21
22 325082      0
23 76      3
24 1      1
25      229

```

1 **Exercise:** Inspect the output from

```
> fulldata[!complete.cases(fulldata), ]
```

2 What does this say about the cases with missing values?

3 _____

4 _____

5 _____

6 _____

7 _____

8 _____

9 _____

10

- 1 Of course, more sophisticated visualization of the cases with missing data
- 2 would likely provide more information as to the source of the missingness.
- 3 In the next Part we will consider such tools.